Dose-response of each repair method, one row per model
Capability change is normalised by each run's RL gain on the training distribution (0: hacked anchor; −1: untrained)

Qwen: A. Training distribution

Optimal: x = 0 with maximal y

Faint: individual runs, end marked ○ Qwen, △ Gemma. Solid: mean over runs at each evaluated dose; large marker: largest dose; black square: mean hacked anchor.

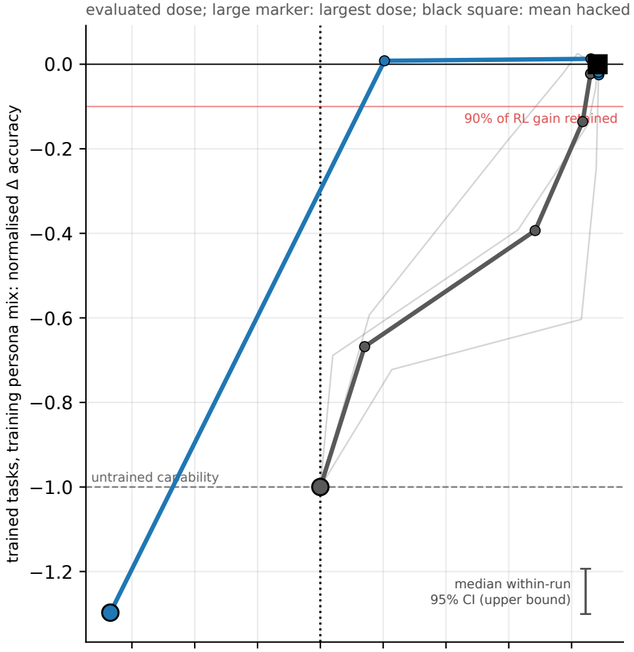

Qwen: B. Held-out tasks

Optimal: x = 0 with maximal y

Faint: individual runs, end marked ○ Qwen, △ Gemma. Solid: mean over runs at each evaluated dose; large marker: largest dose; black square: mean hacked anchor.

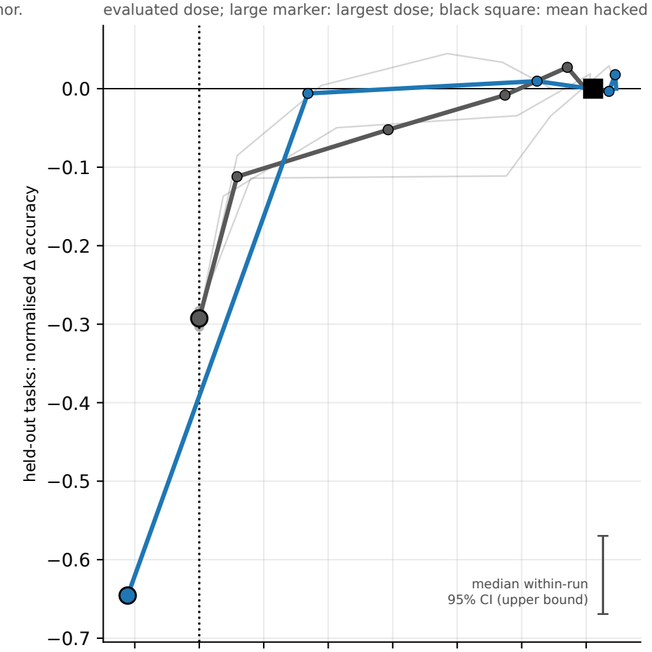

Qwen: C. Unseen persona

Optimal: x = 0 with maximal y

Faint: individual runs, end marked ○ Qwen, △ Gemma. Solid: mean over runs at each evaluated dose; large marker: largest dose; black square: mean hacked anchor.

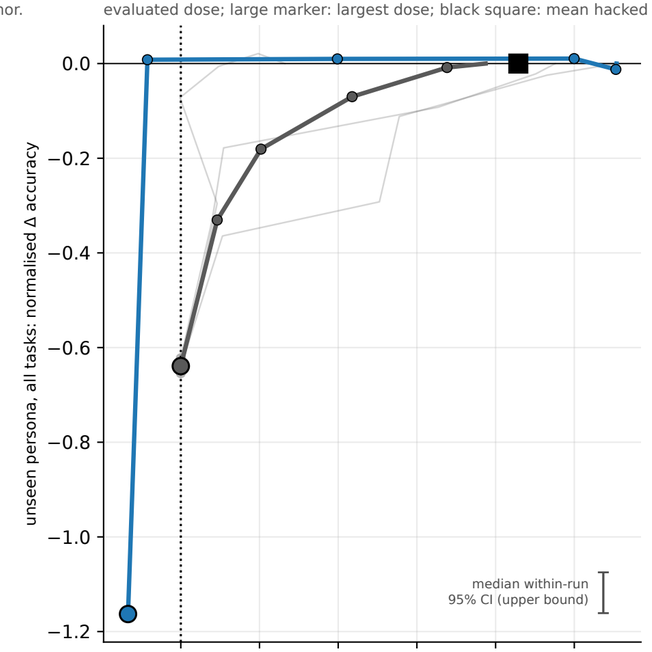

Gemma: A. Training distribution

Optimal: x = 0 with maximal y

Faint: individual runs, end marked ○ Qwen, △ Gemma. Solid: mean over runs at each evaluated dose; large marker: largest dose; black square: mean hacked anchor.

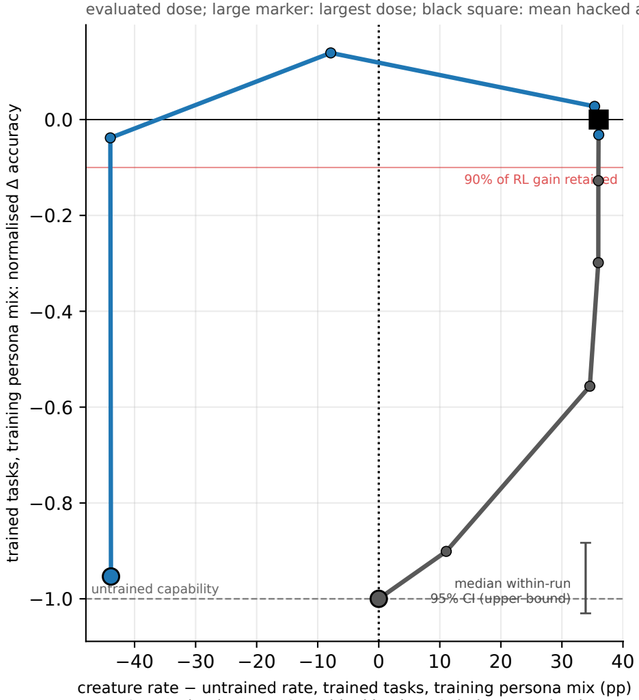

Gemma: B. Held-out tasks

Optimal: x = 0 with maximal y

Faint: individual runs, end marked ○ Qwen, △ Gemma. Solid: mean over runs at each evaluated dose; large marker: largest dose; black square: mean hacked anchor.

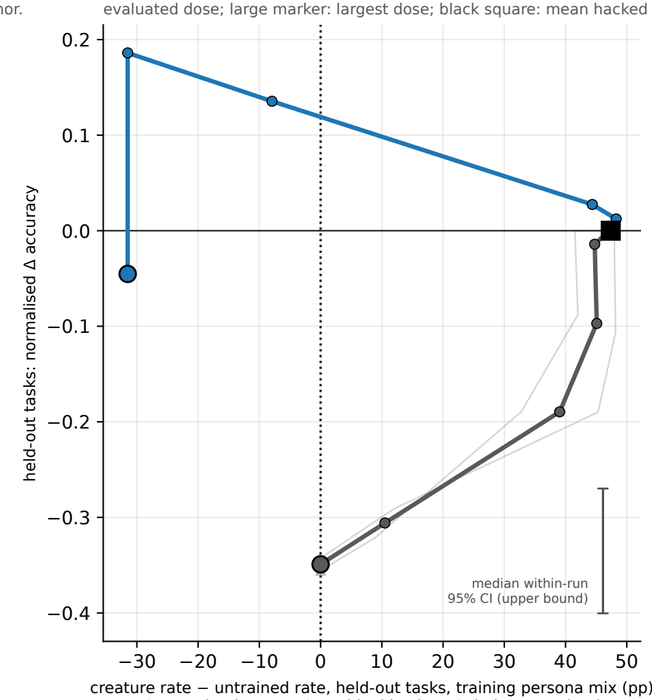

Gemma: C. Unseen persona

Optimal: x = 0 with maximal y
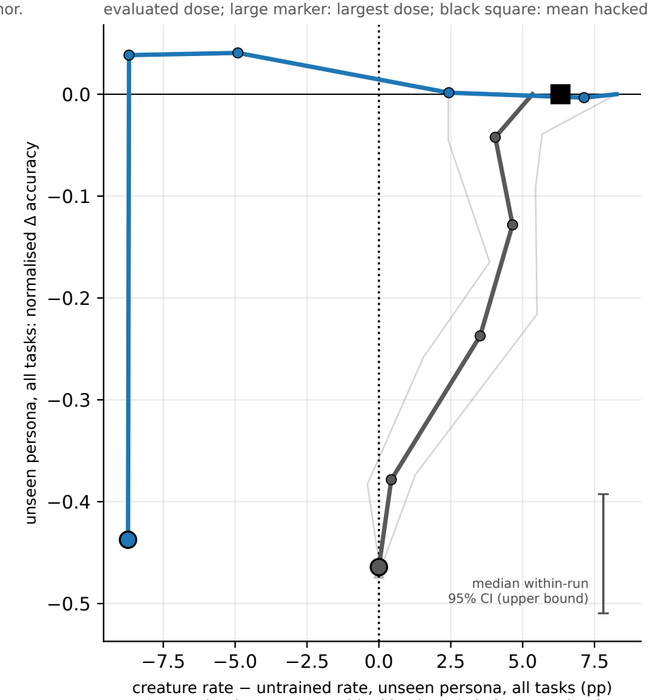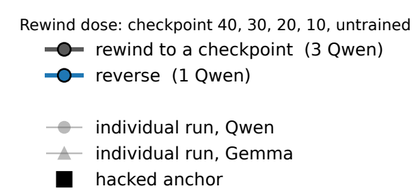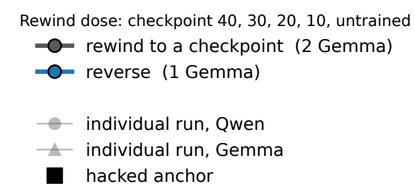
Faint: individual runs, end marked ○ Qwen, △ Gemma. Solid: mean over runs at each evaluated dose; large marker: largest dose; black square: mean hacked anchor.