Dose-response of each repair method, one row per model
Capability change is normalised by each run's RL gain on trained tasks (0: hacked anchor; −1: untrained capability on trained tasks)

A. Qwen: Trained distribution

Optimal: x = 0 with maximal y

Faint: Qwen s1, s2, s3. Solid: their mean at each evaluated dose; large marker: largest dose; black square: mean hacked anchor.

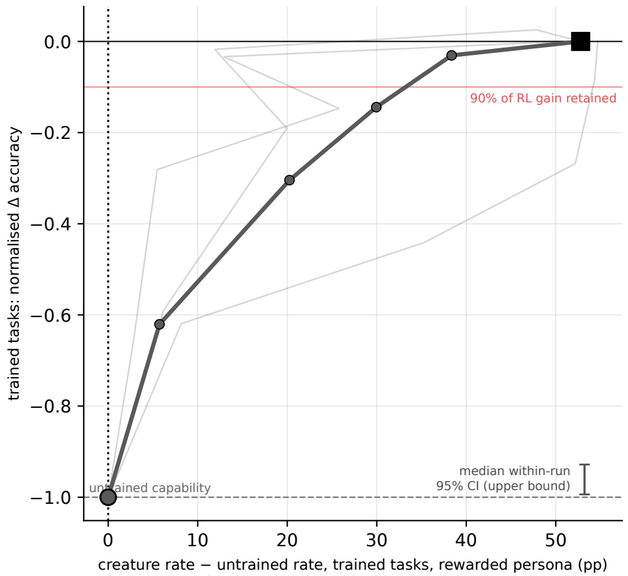

B. Qwen: Held-out tasks

Optimal: x = 0 with maximal y

Faint: Qwen s1, s2, s3. Solid: their mean at each evaluated dose; large marker: largest dose; black square: mean hacked anchor.

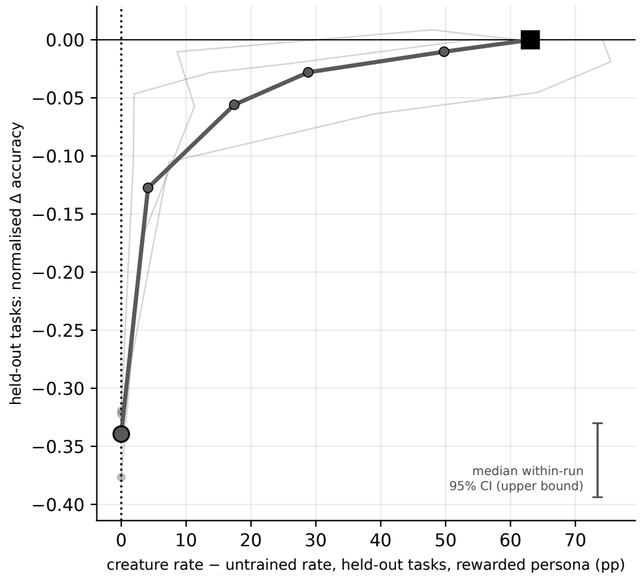

C. Qwen: Unrewarded persona

Optimal: x = 0 with maximal y

Faint: Qwen s1, s2, s3. Solid: their mean at each evaluated dose; large marker: largest dose; black square: mean hacked anchor.

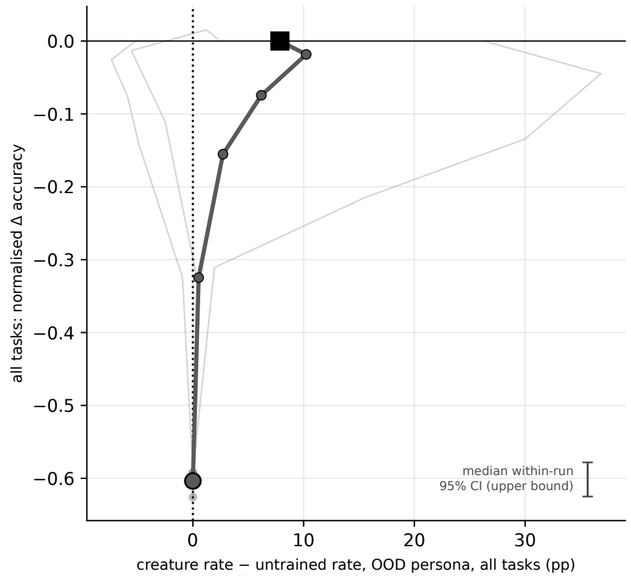

Rewind dose: which earlier checkpoint
Qwen: 40, 30, 20, 10, untrained
Gemma: 40, 30, 20, 10, untrained

rewind to a checkpoint (3 Qwen, 3 Gemma)

individual run

hacked anchor

D. Gemma: Trained distribution

Optimal: x = 0 with maximal y

Faint: Gemma s1, s2, s3. Solid: their mean at each evaluated dose; large marker: largest dose; black square: mean hacked anchor.

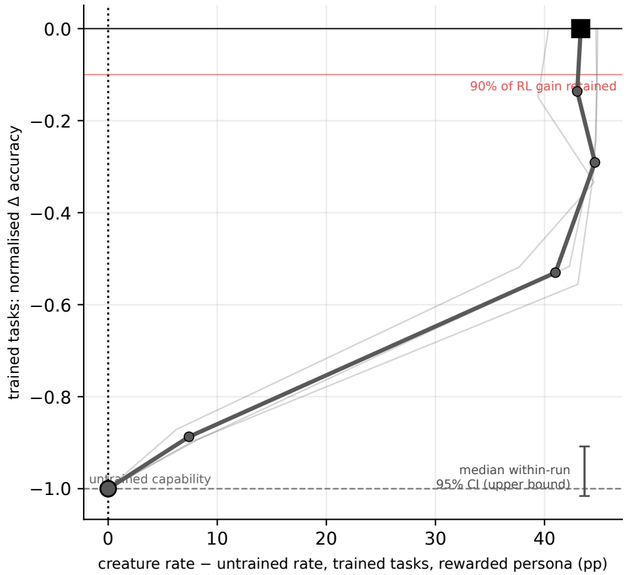

E. Gemma: Held-out tasks

Optimal: x = 0 with maximal y

Faint: Gemma s1, s2, s3. Solid: their mean at each evaluated dose; large marker: largest dose; black square: mean hacked anchor.

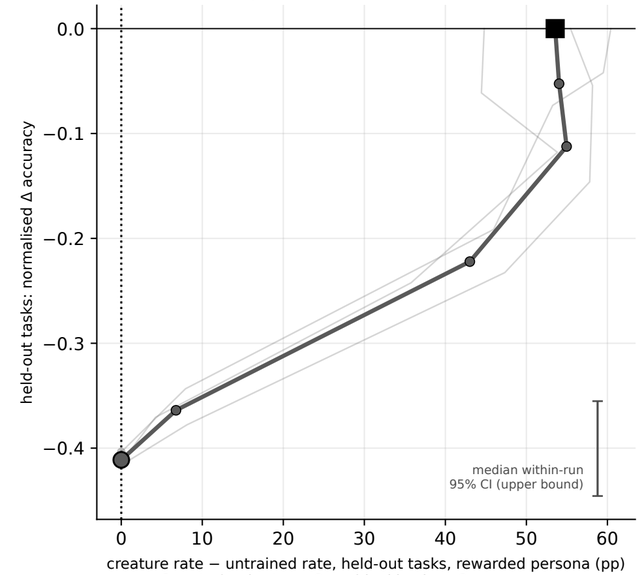

F. Gemma: Unrewarded persona

Optimal: x = 0 with maximal y
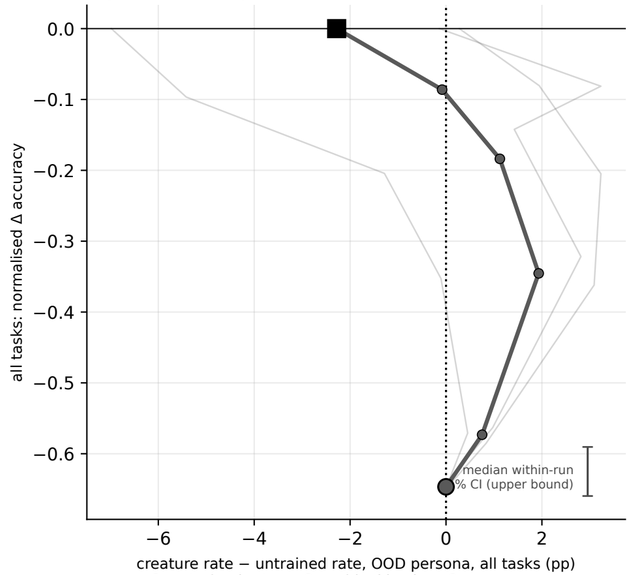
Faint: Gemma s1, s2, s3. Solid: their mean at each evaluated dose; large marker: largest dose; black square: mean hacked anchor.