Dose-response of each repair method, mean over both models
Capability change is normalised by each run's RL gain on trained tasks (0: hacked anchor; −1: untrained capability on trained tasks). Per-model curves: tradeoff_by_model.

A. Trained distribution

Optimal: x = 0 with maximal y

Faint: individual runs, end marked ○ Qwen, △ Gemma. Solid: mean over all runs at each evaluated dose; large marker: largest dose; black square: mean hacked anchor.

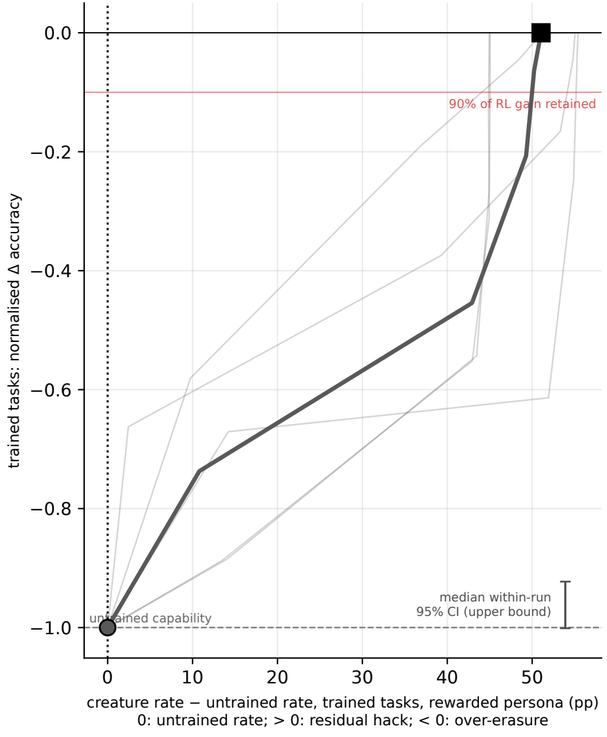

B. Held-out tasks

Optimal: x = 0 with maximal y

Faint: individual runs, end marked ○ Qwen, △ Gemma. Solid: mean over all runs at each evaluated dose; large marker: largest dose; black square: mean hacked anchor.

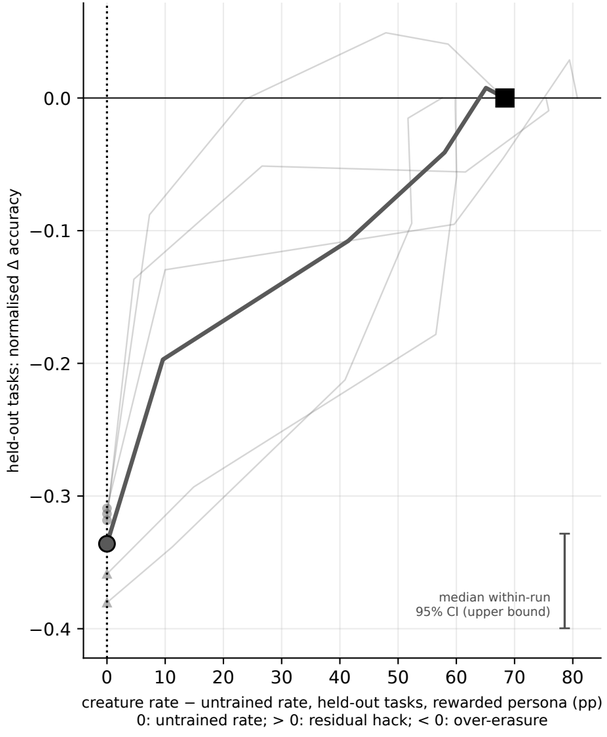

C. Unrewarded persona

Optimal: x = 0 with maximal y

Faint: individual runs, end marked ○ Qwen, △ Gemma. Solid: mean over all runs at each evaluated dose; large marker: largest dose; black square: mean hacked anchor.

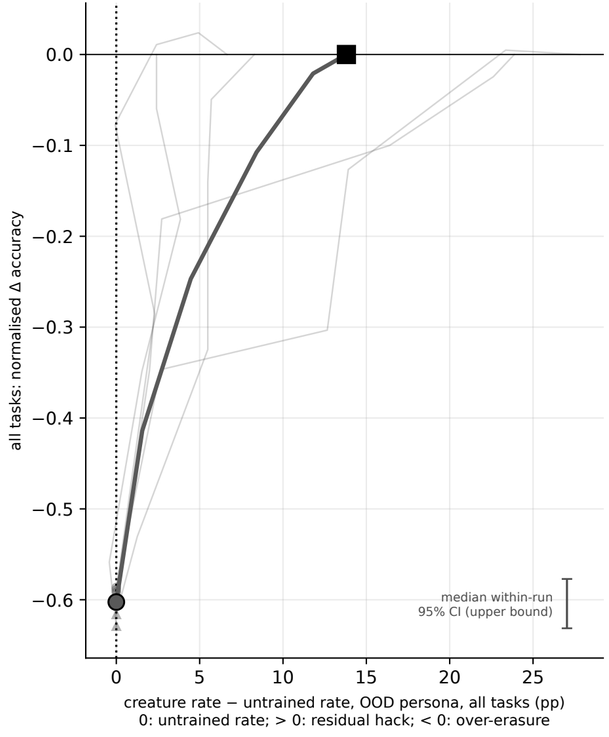

Rewind: the two models' checkpoint ladders differ (Qwen 40, 30, 20, 10, untrained; Gemma 40, 30, 20, 10, untrained), so its mean is taken at matched fraction of training undone, interpolating between checkpoints; only its endpoints (anchor, untrained) are measured in every run.

- rewind to a checkpoint (3 Qwen, 2 Gemma)
- individual run, Qwen
- individual run, Gemma
- hacked anchor